

SDCEL

SCALABLE OVERLAY OPERATIONS OVER DCEL POLYGON LAYERS

Andres Calderon · acald013@ucr.edu

Amr Magdy · amr@cs.ucr.edu


Vassilis J. Tsotras · tsotras@cs.ucr.edu

University of California, Riverside


August 24, 2023

WHAT IS A DCEL?

- Doubly Connected Edge List - DCEL.
- A spatial data structure collecting topological and geometric information for vertices, edges and faces contained by a surface in the plane.
- Widely used to support polygon triangulation and its applications (art gallery problem, robot motion planing, circuit board printing, etc).



Planar subdivision



DCEL representation

DCEL DESCRIPTION

- DCEL uses three tables: Vertices, Faces and Half-edges.




Table 1: **Vertex records.**

vertex	coordinates	incident edge
a	(0,2)	\vec{ba}
b	(2,0)	\vec{db}
⋮	⋮	⋮
e	(4,6)	\vec{fe}
f	(6,4)	\vec{df}
⋮	⋮	⋮

Table 2: **Face records.**


face	boundary edge	hole list
f1	\vec{ab}	nil
f2	\vec{fe}	nil
f3	nil	nil

Table 3: **Half-edge records.**

half-edge	origin	face	twin	next	prev
\vec{fe}	f	f2	\vec{ef}	\vec{ec}	\vec{df}
\vec{ca}	c	f1	\vec{ac}	\vec{ab}	\vec{dc}
\vec{db}	d	f3	\vec{bd}	\vec{ba}	\vec{fd}
⋮	⋮	⋮	⋮	⋮	⋮


ADVANTAGES, CHALLENGES AND CONTRIBUTIONS

- Very efficient for computation of *overlay operators*.
- Allows multiple operations over the same DCEL.
- The output of a DCEL operator can be input to another DCEL operator.




ADVANTAGES, CHALLENGES AND CONTRIBUTIONS

- Very efficient for computation of *overlay operators*.
- Allows multiple operations over the same DCEL.
- The output of a DCEL operator can be input to another DCEL operator.




ADVANTAGES, CHALLENGES AND CONTRIBUTIONS

- Very efficient for computation of *overlay operators*.
- Allows multiple operations over the same DCEL.
- The output of a DCEL operator can be input to another DCEL operator.



ADVANTAGES, CHALLENGES AND CONTRIBUTIONS


- Very efficient for computation of *overlay operators*.
- Allows multiple operations over the same DCEL.
- The output of a DCEL operator can be input to another DCEL operator.



- Currently only sequential DCEL implementations exist.
- Unable to deal with large datasets (i.e. US Census tracks at national level).
- We propose a *scalable* and *distributed* approach to compute the overlay between two DCEL layers.
- Distribution enables scalability, but it also creates challenges: the **orphan-cell** problem and the **orphan-hole** problem.
- We also present *optimizations* that improve the overlay computation performance.


SEQUENTIAL IMPLEMENTATION

- Consider two (simple) input DCELs A_1 and B_1 . The sequential algorithm first finds the intersections of half-edges.
- Then, new vertices (e.g. c_1, c_2) are created, half-edges are updated, new faces are added and labeled (e.g. A_1B_1).




SCALABLE IMPLEMENTATION

- Partition Strategy: we partition the space into cells using a spatial index (e.g. quadtree)
- Each input DCEL layer (e.g. A, B) is partitioned using the index





SCALABLE IMPLEMENTATION

- Each cell should contain all information needed so that it can compute the overlay DCEL locally
- For each cell to be independent, we need to create "artificial" edges and vertices




DISTRIBUTED DCEL CONSTRUCTION



OVERLAY EVALUATION


■ Answering global overlay queries...

- ▶ To compute a particular overlay operator, we query local DCELS.
- ▶ This work is done independently at each cell (node).
- ▶ SDCEL then collects back all local DCEL answers and computes the final answer (by removing artificial edges and concatenating the resulting faces).




LABELING ORPHAN CELLS AND ORPHAN HOLES

- We next discuss the **orphan cell** problem (orphan holes are handled similarly).
- A large face (e.g. the red polygon in the figure) can contain cells that do not intersect with any of the face's boundary edges (called *regular edges*).
- Such cells do not contain any label and thus we do not know which face they belong to.



LABELING ORPHAN CELLS AND ORPHAN HOLES

- We provide an algorithm to efficiently solve the orphan cell problem.



LABELING ORPHAN CELLS AND ORPHAN HOLES

Algorithm 1: GETNEXTCELLWITHEDGES algorithm

Input: a quadtree Q and a list of cells M .

```
1 function GETNEXTCELLWITHEDGES( $Q, M$ ):  
2    $C \leftarrow$  orphan cells in  $M$   
3   foreach orphanCell in  $C$  do  
4     initialize cellList with orphanCell  
5     nextCellWithEdges  $\leftarrow$  nil  
6     referenceCorner  $\leftarrow$  nil  
7     done  $\leftarrow$  false  
8     while  $\neg$ done do  
9        $c \leftarrow$  last cell in cellList  
10      cells, corner  $\leftarrow$  GETCELLSATCORNER( $Q, c$ )  
11      foreach cell in cells do  
12        nedges  $\leftarrow$  get edge count of cell in  $M$   
13        if nedges > 0 then  
14          nextCellWithEdges  $\leftarrow$  cell  
15          referenceCorner  $\leftarrow$  corner  
16          done  $\leftarrow$  true  
17        else  
18          add cell to cellList  
19        end  
20      end  
21    end  
22    foreach cell in cellList do  
23      output(cell, nextCellWithEdges,  
24        referenceCorner)  
25      remove cell from  $C$   
26    end  
27  end
```

Algorithm 2: GETCELLSATCORNER algorithm

Input: a quadtree with cell envelopes Q and a cell c .

```
1 function GETCELLSATCORNER( $Q, c$ ):  
2   region  $\leftarrow$  quadrant region of  $c$  in  $c$ .parent  
3   switch region do  
4     case 'SW' do  
5       corner  $\leftarrow$  left bottom corner of  $c$ .envelope  
6     case 'SE' do  
7       corner  $\leftarrow$  right bottom corner of  $c$ .envelope  
8     case 'NW' do  
9       corner  $\leftarrow$  left upper corner of  $c$ .envelope  
10    case 'NE' do  
11      corner  $\leftarrow$  right upper corner of  $c$ .envelope  
12  end  
13  cells  $\leftarrow$  cells which intersect corner in  $Q$   
14  cells  $\leftarrow$  cells -  $c$   
15  cells  $\leftarrow$  sort cells on basis of their depth  
16  return (cells, corner)  
17 end
```

- Optimizing for faces overlapping many cells...
 - ▶ Naive approach sends all faces that overlap a cell to a master node (that will combine them).
 - ▶ We propose an intermediate reduce processing step.
 - The user provides a level in the quadtree structure and faces are evaluated at those intermediate reducers.
 - ▶ We also consider another approach that re-partitions such faces using their labels as the key.
 - It avoids the reduce phase but implies an additional shuffle.
 - However, as we show in the experiments this overhead is minimal.

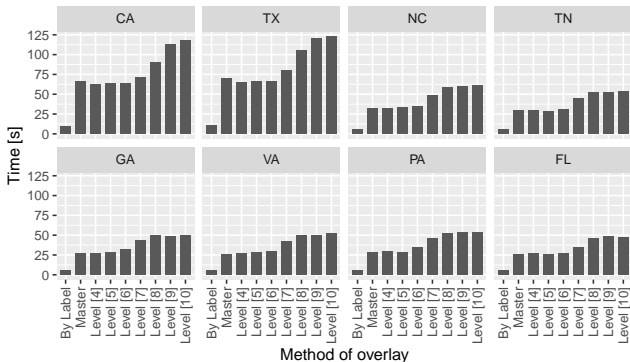
- Optimizing for unbalanced layers...
 - ▶ Finding intersections is the most critical part of the overlay computation.
 - ▶ However, in many cases one of the layers has much more half-edges than the other.
 - ▶ Sweep-line algorithms to detect intersections run over all the edges.
 - ▶ Instead we scan the larger dataset only for the x-intervals where there are half-edges from the smaller dataset.

■ Datasets.

Dataset	Layer	Number of polygons	Number of edges
MainUS	Polygons for 2000	64983	35417146
	Polygons for 2010	72521	36764043
GADM	Polygons for Level 2	116995	32789444
	Polygons for Level 3	117891	37690256
CCT	Polygons for 2000	7028	2711639
	Polygons for 2010	8047	2917450



EXPERIMENTAL EVALUATION

■ Evaluation of the overlapping faces optimization.




EXPERIMENTAL EVALUATION

■ Evaluation of the unbalanced layers optimization.





- Performance varying number of partition cells (CCT dataset).





EXPERIMENTAL EVALUATION

■ Performance with MainUS and GADM datasets.





EXPERIMENTAL EVALUATION

■ MainUS scale-up and speed-up.



■ GADM scale-up and speed-up.



- We introduced SDCEL, a scalable approach to compute the overlay operation among two layers that represent polygons from a planar subdivision of a surface.
- We use a partition strategy which guarantees that each partition (cell) has the data needed to work independently.
- We also proposed several optimizations to improve performance.
- Our experiments using real datasets show very good performance; we are able to compute overlays over very large layers (each with >35M edges) in few minutes.

Thank you!