

SDCEL

SCALABLE OVERLAY OPERATIONS OVER DCEL POLYGON LAYERS

Andres Calderon · acald013@ucr.edu

Amr Magdy · amr@cs.ucr.edu


Vassilis Tsotras · tsotras@cs.ucr.edu

University of California, Riverside

August 22, 2023

WHAT IS A DCEL?

- Doubly Connected Edge List - DCEL.
- A spatial data structure collecting topological and geometric information for vertices, edges and faces contained by a surface in the plane.
- Widely use to support polygon triangulation and its applications (art gallery problem, robot motion planing, circuit board printing, etc).



DCEL DESCRIPTION




Table 1: Vertex records.

| vertex | coordinates | incident edge |
|----------|-------------|---------------|
| a | (0,2) | \vec{ba} |
| b | (2,0) | \vec{db} |
| c | (2,4) | \vec{dc} |
| \vdots | \vdots | \vdots |

Table 2: Face records.


| face | boundary edge | hole list |
|-------|---------------|-----------|
| f_1 | \vec{ab} | nil |
| f_2 | \vec{fe} | nil |
| f_3 | nil | nil |

Table 3: Half-edge records.


| half-edge | origin | face | twin | next | prev |
|------------|----------|----------|------------|------------|------------|
| \vec{fe} | f | f_2 | \vec{ef} | \vec{ec} | \vec{df} |
| \vec{ea} | e | f_1 | \vec{ac} | \vec{ab} | \vec{dc} |
| \vec{db} | d | f_3 | \vec{bd} | \vec{ba} | \vec{fd} |
| \vdots | \vdots | \vdots | \vdots | \vdots | \vdots |

ADVANTAGES, CHALLENGES AND CONTRIBUTIONS


- Very efficient for computation of *overlay maps* and *overlay operators*.
- Allows multiple operations over the same DCEL and chaining.
- Currently only sequential implementations.
- Unable to deal with large datasets (i.e. US Census tracks at national level).
- We propose a scalable and distributed approach to compute the overlay between two DCEL layers.
- It solve issues related to holes and large empty areas in addition to optimizations during the overlay computations.





SEQUENTIAL IMPLEMENTATION



PARTITION STRATEGY




DISTRIBUTED DCEL CONSTRUCTION




LABELING ORPHAN CELLS AND HOLES

- The **orphan cell** problem.
- When a cell is empty (it does not intersect or contain a *regular edge*).
- A regular edge is one that is not part of a hole.
- The label for the cell (or a hole if it is present) is lost. Thus we do not know which face it belongs to.
- We provide algorithms to solve the issue.



LABELING ORPHAN CELLS AND HOLES



LABELING ORPHAN CELLS AND HOLES

Algorithm 1: GETNEXTCELLWITHEDGES algorithm

Input: a quadtree Q and a list of cells M .

```
1 function GETNEXTCELLWITHEDGES( $Q, M$ ):  
2    $C \leftarrow$  orphan cells in  $M$   
3   foreach orphanCell in  $C$  do  
4     initialize cellList with orphanCell  
5     nextCellWithEdges  $\leftarrow$  nil  
6     referenceCorner  $\leftarrow$  nil  
7     done  $\leftarrow$  false  
8     while  $\neg$ done do  
9        $c \leftarrow$  last cell in cellList  
10      cells, corner  $\leftarrow$  GETCELLSATCORNER( $Q, c$ )  
11      foreach cell in cells do  
12        nedges  $\leftarrow$  get edge count of cell in  $M$   
13        if nedges > 0 then  
14          nextCellWithEdges  $\leftarrow$  cell  
15          referenceCorner  $\leftarrow$  corner  
16          done  $\leftarrow$  true  
17        else  
18          add cell to cellList  
19        end  
20      end  
21    end  
22    foreach cell in cellList do  
23      output(cell, nextCellWithEdges,  
24        referenceCorner)  
25      remove cell from  $C$   
26    end  
27  end
```

Algorithm 2: GETCELLSATCORNER algorithm


Input: a quadtree with cell envelopes Q and a cell c .

```
1 function GETCELLSATCORNER( $Q, c$ ):  
2   region  $\leftarrow$  quadrant region of  $c$  in  $c$ .parent  
3   switch region do  
4     case 'SW' do  
5       corner  $\leftarrow$  left bottom corner of  $c$ .envelope  
6     case 'SE' do  
7       corner  $\leftarrow$  right bottom corner of  $c$ .envelope  
8     case 'NW' do  
9       corner  $\leftarrow$  left upper corner of  $c$ .envelope  
10    case 'NE' do  
11      corner  $\leftarrow$  right upper corner of  $c$ .envelope  
12  end  
13  cells  $\leftarrow$  cells which intersect corner in  $Q$   
14  cells  $\leftarrow$  cells -  $c$   
15  cells  $\leftarrow$  sort cells on basis of their depth  
16  return (cells, corner)  
17 end
```

OVERLAY EVALUATION

■ Answering global overlay queries...

- ▶ We query local DCEL's for particular overlay operators.
- ▶ It filters out at each worker.
- ▶ SDCEL collects back and reduce the final answer (remove artificial edges and concatenate splits).



- Optimizations for faces expanding cells...
 - ▶ Naive approach send all faces to a master node.
 - ▶ We propose an intermediate reduce processing step.
 - ▶ The user provides a level in the quadtree structure and faces are evaluated at those intermediate reducers.
 - ▶ Another approach re-partitions the faces using its labels as the key.
 - ▶ It avoids the reduce phase but implies an additional shuffle.
 - ▶ However, SDCEL works with much smaller independent amounts of work.

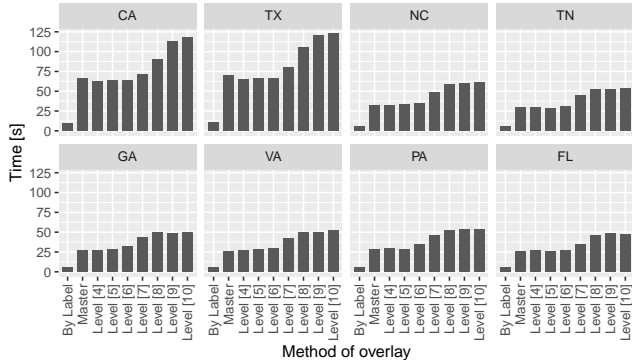
- Optimizing for unbalanced layers...
 - ▶ Finding intersections is the most critical part of the overlay computation.
 - ▶ However, in many cases one of the layer has much more half-edges than the other.
 - ▶ Sweep-line algorithms to detect intersections runs over all the edges.
 - ▶ We scan the larger dataset only for the x-intervals where there are half-edges for the smaller dataset.
 - ▶ It avoids unnecessary scanning where there are few edge (or even none) from one of the layers.

■ Datasets.

| Dataset | Layer | Number of polygons | Number of edges |
|---------|----------------------|-----------------------|--------------------|
| MainUS | Polygons for 2000 | 64983 | 35417146 |
| | Polygons for 2010 | 72521 | 36764043 |
| GADM | Polygons for Level 2 | 116995 | 32789444 |
| | Polygons for Level 3 | 117891 | 37690256 |
| CCT | Polygons for 2000 | 7028 | 2711639 |
| | Polygons for 2010 | 8047 | 2917450 |



EXPERIMENTAL EVALUATION

■ Evaluation of overlay methods.





EXPERIMENTAL EVALUATION

■ Unbalance layers optimization.





EXPERIMENTAL EVALUATION

- Performance varying number of cells (CCT dataset).



EXPERIMENTAL EVALUATION

■ Performance with MainUS and GADM datasets.



- We introduced SDCEL, a scalable approach to compute the overlay operation among two layers that represent polygons from a planar subdivision of a surface.
- We first presented a partition strategy which guarantee each partition has the needed data to work independently.
- We also proposed several optimization to improve performance and ensure correct results.
- Our experiments in real datasets show very good performance and it is able to compute overlays over very large layers in few minutes.